

There is a subtle but important shift taking place in the relationship between executives and artificial intelligence. The question is no longer simply whether AI can help executives make better decisions. Increasingly, the question is whether the executive is still making the judgement at all.

Consider a CEO preparing for a board presentation. The organization is considering entering a new market, and the CEO needs to explain the opportunity, defend the financial case, anticipate the risks and persuade the board that the organization should commit substantial resources to the move. Traditionally, the executive would gather the evidence, test the assumptions, consult colleagues, weigh competing interests and eventually form a judgement. The board might challenge the recommendation, but the CEO remains the person who has interpreted the situation and decided what ought to be done.

Now introduce an AI model. The CEO can provide the model with financial statements, market intelligence, competitor information, customer data, strategic plans and previous board papers. Within minutes, it can identify patterns, generate scenarios, challenge assumptions, construct the presentation and recommend a course of action. It can even anticipate the questions individual board members are likely to ask. On the surface, this looks like a significant improvement in executive decision-making.

But something more fundamental has happened. The CEO is no longer simply using AI to analyse a decision. The model has become part of the process through which the CEO *sees* the decision. The

question therefore becomes less about whether the model is correct and more about what happens when its interpretation enters the executive's own process of judgement.

When the model gives you the answer, who is making the judgement?

The model changes the journey from information to judgement

Information does not become judgement simply because it has been processed more quickly or intelligently. The model may establish that the new market has attractive growth prospects, that the projected return exceeds the company's investment hurdle rate and that the probability-adjusted downside appears manageable. But none of these calculations, by themselves, determines what the organization *ought* to do.

The CEO still has to decide what matters. Is growth more important than organizational stability? Is the projected return sufficient compensation for geopolitical risk? Does the organization have an obligation to protect existing employees before pursuing expansion? Is entering this market consistent with what the organization believes itself to be? These are questions of executive judgement, and executive judgement is shaped by something deeper than the information available to the executive.

It is shaped by worldview.

The worldview beneath the decision

The **Worldview LOA Framework** provides the theoretical foundation for understanding this deeper layer of executive judgement. It conceptualizes executive worldviews through **Legitimacy, Obligation and Agency**, showing how these deep evaluative structures shape perception, judgement and action through **Evaluative Reflexivity**.

The CEO therefore encounters the model's recommendation through an existing evaluative position. **Legitimacy** concerns what the executive considers right, credible, appropriate or worthy of trust. **Obligation** concerns what the executive believes ought to be done and the responsibilities and accountabilities that accompany their position. **Agency** concerns what the executive believes they can influence, mobilise or change.

These questions do not disappear when AI becomes more capable. They become more important. The model might say, *"This is the optimal strategy,"* but the executive still needs to ask: optimal according to what, for whom, based on which assumptions, consistent with which obligations and within whose perceived capacity to act?

The model can generate the recommendation. It cannot determine the evaluative position from which that recommendation will ultimately be judged.

The banker faces the same problem

Consider a private banker advising a high-net-worth investor. The investor asks whether she should move a substantial portion of her wealth into a particular asset class. The banker now has access to a model capable of analysing thousands of market variables, constructing scenarios, comparing portfolios and generating sophisticated risk-return assessments.

The model recommends Portfolio A, perhaps because it offers the highest risk-adjusted return. The banker could simply present that recommendation to the client. But the client may be approaching retirement, her wealth may have been accumulated over generations, and she may be deeply uncomfortable with permanent capital loss. She may want to preserve a family business, fund her children's education or establish a philanthropic foundation. The model can optimise the portfolio, but it cannot, by itself, determine what the wealth *means* to the person who owns it.

The client asks, *"What should I do with my wealth?"* The model may silently translate that into, *"Which portfolio maximises risk-adjusted return within these parameters?"* Those are not the same question. The first is a question of human judgement; the second is an optimisation problem. The model may be extraordinarily good at solving the second while never establishing whether it is actually the right problem to solve.

Take the Fix

This is where the **Executive Sextant**™ becomes important. The Worldview LOA Framework establishes an executive's evaluative "**fix**" through Legitimacy, Obligation and Agency. Through Evaluative Reflexivity, the executive brings these three bearings together, examines where they align or conflict, and establishes a reliable evaluative position. The Executive Sextant then provides the process through which that position is applied to executive judgement.

The metaphor matters. A navigator does not set a course before establishing a position. Executive judgement should be no different. In an AI-enabled environment, the Executive Sextant creates a deliberate space between **model output and executive action**. Its six stages are **Observation, Appraisal, Angle of Appraisal, Judgement, Engagement and Behaviour**.

Observation

The first stage is **Observation**: attending to what is actually present in the environment, relationships and inner landscape without immediately interpreting or judging it. AI can dramatically expand what an executive is able to observe, but the executive must distinguish between what is actually present and what the model has already interpreted. The CEO therefore needs to ask what is happening before accepting the model's framing of the situation. The banker needs to distinguish the client's actual circumstances from the financial representation generated by the model.

Appraisal

The second stage is **Appraisal**, the first-pass evaluation of what has been observed through the executive's existing schema. This makes explicit the automatic assessments that typically operate below conscious awareness. The CEO may see an attractive growth opportunity where another executive sees unacceptable strategic risk. The banker may see an opportunity to improve portfolio performance where another sees a threat to intergenerational wealth. The data may be similar, but the appraisal is not.

Angle of Appraisal

The third stage, **Angle of Appraisal**, is the critical diagnostic point. It requires the executive to examine the precise angle from which they are viewing the situation, recognising that small shifts in that angle can produce dramatically different interpretations and responses. This is where the LOA bearings become particularly useful. What is legitimate? What am I obligated to do? What can I actually do?

The same discipline should be applied to the model itself. AI recommendations can appear objective, but every model operates within assumptions about what matters, which variables are relevant and what constitutes a desirable outcome. The executive therefore needs to recognise not only their own angle of appraisal, but also the angle embedded in the model's framing of the problem.

Judgement

The fourth stage is **Judgement**: the considered verdict that integrates observation, appraisal and perspective-awareness into a coherent assessment of the situation. The CEO may accept the model's recommendation, modify it or reject it. The important point is that the executive knows why.

The banker may conclude that Portfolio B is more appropriate because stewardship and preservation of capital outweigh incremental expected return. Neither executive is rejecting AI. They are exercising judgement *after* AI, rather than allowing the model's conclusion to become judgement by default.

Engagement

The fifth stage is **Engagement**: how the executive chooses to enter the situation, including the quality, energy and orientation of their relational and strategic engagement. The CEO can present the model's recommendation to the board as an answer, or as a proposition that deserves challenge. The banker can say, *"The model recommends this portfolio,"* or instead say, *"The analysis suggests this option, but let's examine whether it actually fits what you are trying to protect."*

The difference is not merely semantic. It reflects whether the executive is using AI to replace judgement or to deepen the conversation in which judgement is exercised.

Behaviour

The sixth stage is **Behaviour**, the visible manifestation of the full diagnostic cycle through the decisions, actions and communications that flow from worldview through engagement. This is where judgement becomes consequential. The CEO commits capital. The banker recommends the portfolio. The organization acts. The client invests.

AI may have contributed enormously to the quality and speed of the analysis, but the resulting behaviour remains attributable to the human decision-maker.

The model can accelerate the voyage. It cannot take the fix.

The temptation in the AI debate is to ask whether machines will eventually become better decision-makers than humans. A more important question may be whether humans will remain sufficiently aware of the process through which they turn information into judgement.

AI can generate possibilities, expose patterns, model consequences and challenge assumptions. But the Executive Sextant asks what happens **between the model's answer and human action**. That space is where executive judgement survives.

The Worldview LOA Framework provides the deeper evaluative bearings of **Legitimacy, Obligation and Agency**, while the Executive Sextant provides the process through which those bearings are brought into the movement from **Observation to Behaviour**. Together, they offer a discipline for AI-enabled leadership: not rejecting the model, but refusing to confuse its answer with one's own judgement.

The CEO remains accountable to the board. The banker remains accountable to the client. The executive remains accountable for the consequences of action.

So perhaps the question is not simply:

> **Can AI make better decisions than executives?**

It is:

> **When the model gives you the answer, who is making the judgement?**

A navigator never sets a course before taking a fix.

Executive judgement should be no different.